

Zeyi Liao

☎ 213-551-8327 | ✉ lzy37ld@gmail.com / liao.629@osu.edu | 📱 X | 🎓 Google Scholar | 🌐 Website

🔌 Research Interests

I am broadly interested in AI agents, aiming to build helpful *yet aligned* agent systems that benefit humanity. AI models have crossed a critical capability threshold for general-purpose work; my next mission is to harness the rogue AI and build trustworthy autonomy.

🎓 Education

- **The Ohio State University** (Advisor: **Huan Sun**) Aug. 2023 – Present
Ph.D. in Computer Science and Engineering
- **University of Southern California** (Advisor: **Xiang Ren**) Aug. 2021 – May. 2023
M.S in Electronic and Computer Engineering
- **Beijing Jiaotong University** Aug. 2016 – May. 2020
B.S in Electronic Engineering

🧩 Experience

- **Research Intern**, Neocognition Apr. 2026 – Present
Host: **Xiang Deng**
Topic: Continual learning, agent harness, long-horizon computer-use agents
- **Research Intern**, Microsoft Research (Redmond) May. 2025 – Aug. 2025
Host: **Ahmed Awadallah, Yadong Lu**
Topic: Computer-use agents, visual grounding

📖 Publications (selected, see full list on 🎓)

* means equal contribution

- **RedTeamCUA: Realistic Adversarial Testing of Computer-Use Agents in Hybrid Web-OS Environments**
Zeyi Liao*, Jaylen Jones*, Linxi Jiang*, Eric Fosler-Lussier, Yu Su, Zhiqiang Lin, Huan Sun
International Conference on Learning Representations (ICLR), 2026. [PDF] **Oral (1.5%)**
‣ The only safety benchmark adopted in the evaluation of **Qwen-CUA**.
- **EIA: Environmental Injection Attack on Generalist Web Agents for Privacy Leakage**
Zeyi Liao*, Lingbo Mo*, Chejian Xu, Mintong Kang, Jiawei Zhang, Chaowei Xiao, Yuan Tian, Bo Li, Huan Sun
International Conference on Learning Representations (ICLR), 2025. [PDF]
‣ Selected as course material for **UCSD CSE 291: LLM Security**.
‣ Incorporated into UC Berkeley RDI's **SuperRed**.
- **Amplegcg: Learning a universal and transferable generative model of adversarial suffixes for jailbreaking both open and closed llms**
Zeyi Liao, Huan Sun
Conference on Language Modeling (COLM), 2024. [PDF]
‣ Discussed in **Daniel Paleka's AI safety newsletter**.
‣ Follow-up work: **AmpleGCG-plus** (🔥 X thread).

- **Agent Learning via Early Experience**

Kai Zhang, Xiangchao Chen*, Bo Liu*, Tianci Xue*, Zeyi Liao*, *et al.*, Huan Sun, Jason Weston, Yu Su, Yifan Wu (30 authors)

International Conference on Machine Learning (ICML), 2026. [\[PDF\]](#)

- **Mind2Web 2: Evaluating Agentic Search with Agent-as-a-Judge**

Boyu Gou, Zanming Huang, Yuting Ning, Yu Gu, Michael Lin, Weijian Qi, *et al.*, Zeyi Liao, *et al.*, Huan Sun, Yu Su (26 authors)

Conference on Neural Information Processing Systems (NeurIPS), 2025. [\[PDF\]](#)

- **Beyond Clicking: A Step Towards Generalist GUI Grounding via Text Dragging**

Zeyi Liao, Yadong Lu, Boyu Gou, Huan Sun, Ahmed Awadallah

Preprint on Arxiv, 2025. [\[PDF\]](#)

- **WebGuard: Building a Generalizable Guardrail for Web Agents**

Boyuan Zheng, Zeyi Liao, Scott Salisbury, Zeyuan Liu, Michael Lin, Qinyuan Zheng, Zifan Wang, Xiang Deng, Dawn Song, Huan Sun, Yu Su

Preprint on Arxiv, 2025. [\[PDF\]](#)

- **AdvAgent: Controllable Blackbox Red-teaming on Web Agents**

Chejian Xu, Mintong Kang, Jiawei Zhang, Zeyi Liao, Lingbo Mo, Mengqi Yuan, Huan Sun, Bo Li

International Conference on Machine Learning (ICML), 2025 [\[PDF\]](#)

- **Joint Demonstration and Preference Learning Improves Policy Alignment with Human Feedback**

Chenliang Li, Siliang Zeng*, Zeyi Liao*, Jiaxiang Li, Dongyeop Kang, Alfredo Garcia, Mingyi Hong

International Conference on Learning Representations (ICLR), 2025. [\[PDF\]](#)

Service

- **Reviewer:** COLM, CVPR, EMNLP, ICLR, IEEE TNNLS, NeurIPS

- **Teaching Assistant:** CSE 3521 Survey of Artificial Intelligence I @ OSU Aug. 2023 – Jan. 2024

Awards

- **Master Honor Student**, University of Southern California

Aug. 2022

- **First Class Scholarship**, Beijing Jiaotong University

Aug. 2020/2019/2018

- **Merit Student**, Beijing Jiaotong University

Aug. 2018/2017